

A thick black L-shaped frame surrounds the central text. The top-left corner has a horizontal bar extending to the right and a vertical bar extending downwards. The bottom-right corner has a horizontal bar extending to the left and a vertical bar extending upwards.

INTER-PROCESS COMMUNICATION

Tanzir Ahmed
CSCE 313 Spring 2021

Inter-Process Communication


■ IPC Methods

- *Pipes and FIFO*
- *Message Passing*
- *Shared Memory*
- *Semaphore Sets*
- *Signals*

■ References:


- *Beej's guide to Inter Process Communication for the code examples* (<https://beej.us/guide/bgipc/>)
- *Understanding Unix/Linux Programming*, Bruce Molay, Chapters 10, 15
- [*Advanced Linux Programming Ch 5*](#)
- *Some material also directly taken or adapted with changes from [Illinois course in System Programming](#) (Prof. Angrave), [UCSD](#) (Prof. Snoeren), and [USNA](#) (Prof. Brown)*

IPC Fundamental Communication Models



(a)

Example: pipe, fifo, message, signal




(b)

Example: shared memory, memory mapped file


Unnamed Pipes

```
int pipe(int fildes[2])
```


- Returns a pair of file descriptors
 - *fildes[0] is the read end and fildes[1] is the write end*
- Create a message pipe
 - *Data is received in the order it was sent*
 - *OS enforces mutual exclusion: only one process at a time*
 - *Processes sharing the pipe must have same parent in common*
 - *The space in between (in Kernel) is bounded (i.e., you cannot send unlimited msgs w/o receiving)*



BEFORE



AFTER



IPC Pipe - Method

```
#include <stdio.h>
#include <unistd.h>

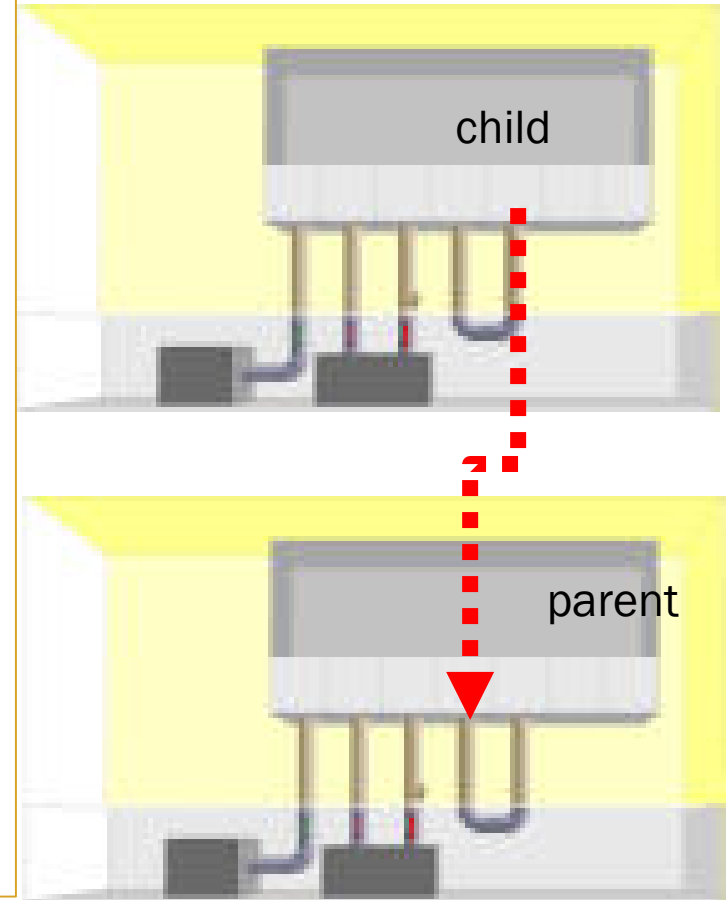
void main ()
{
    char buf [10];
    int fds [2];
    pipe (fds);
    printf ("sending msg: Hi\n");
    write (fds[1], "Hi", 3);
    read (fds[0], buf, 3);
    printf ("Received msg: %s\n", buf);
}
```

Connects the
two fds as pipe

```
compute-linux1 tanzir/code> ./a.out
sending msg: Hi
Received msg: Hi
```


Unnamed Pipe Between Two Processes

```
int main ()
{
    int fds [2];
    pipe (fds); // connect the pipe
    if (!fork()){ // on the child side
        sleep (3);
        char * msg = "a test message";
        printf ("CHILD: Sent %s\n", msg);
        write (fds [1], msg,
strlen(msg)+1);
    }else{
        char buf [100];
        read (fds [0], buf, 100);
        printf ("PRNT: Recvd %s\n", buf);
    }
    return 0;
}
```



Named Pipes (FIFO)

- FIFOs are a mechanism that allow for IPC that's to some degree similar to using regular files
 - *Because you have to **open()** the pipe file to start a communication and then **read()/write()** to/from it*
 - *FIFOs files are persistent in disk, just like regular files.*
- However, it is also very different from using regular files in the sense that Data is never actually written to disk (instead it is stored in buffers in memory) so the overhead of disk I/O (which is huge!) is avoided
 - *Filename is only used for system-wide visibility/scope of the pipe, not for containing data*
- FIFOs are similar to unnamed pipes because:
 - *They are unidirectional: 1 side can only either read or write*
 - *Mechanism is a Kernel-managed bounded queue, just like unnamed pipes*



FIFO - Problems

- The processes need to **agree on a name** ahead of time – how to communicate that??

```
FIFORequestChannel rc ("control", ..) {  
    ...  
    mkfifo ("control", PERMS); // create  
}
```

- **Not concurrency safe** within a process
 - *Like a file used by multiple processes/threads*
 - *Multiple threads writing can cause race condition*

Using FIFO's

- `mkfifo (name)`: create a FIFO
- How do I remove a FIFO: `rm fifoname` or `unlink(fifoname)`
- How do I listen at a FIFO for a connection
 - `open (fifoname, O_RDONLY)`
- How do I open a FIFO in write mode?
 - `open(fifoname, O_WRONLY)`
- Can someone open in both read and write mode?
 - *No. That makes it directional, but efficient*

FIFO DEMO

Writer


```
#define FIFO_NAME "test.txt"
int main(void)
{
    char s[300];
    int num, fd;
    mkfifo(FIFO_NAME, 0666); // create
    printf("Waiting for readers...\n");
    fd = open(FIFO_NAME, O_WRONLY); //open
    if (fd < 0)
        return 0;
    printf("Got a reader--type some stuff\n");
    while (gets(s)) {
        if (!strcmp (s, "quit")) break;
        if ((num = write(fd, s, strlen(s)))== -1)
            perror("write");
        else
            printf("SENDER:wrote %d bytes\n",num);
    }
    //unlink (FIFO_NAME);
    return 0;
}
```

Reader

```
int main(void)
{
    char s[300];
    int num, fd;
    printf("waiting for writers...\n");
    fd = open(FIFO_NAME, O_RDONLY);
    printf("got a writer\n");
    do{
        if ((num = read(fd, s, 300)) == -1)
            perror("read");
        else {
            s[num] = '\0';
            printf("RECV: Read %d bytes:
\\\"%s\\\"\\n", num, s);
        }
    } while (num > 0);
    return 0;
}
```

Message Queues

Direct



Indirect



Added Features of MQ

- Supports Priority of messages, effectively changing the FIFO order of the messages (note pipe/FIFO messages are only FIFO)
- Multiple processes can read from or write into the same message queue – not possible in FIFO
- MQ does not require the ends be simultaneously connected
 - *FIFO requires both processes connected*
- Allows for asynchronous delivery of messages
 - *The recipient process sets up notification of message receipt using `mq_notify()`*



Operations on Message Queues

```
mqd_t mq_open(const char *name, int oflag,  
              mode_t mode, struct mq_attr *attr);
```

```
int mq_send(mqd_t mqdes, const char *msg_ptr,  
            size_t msg_len, unsigned int msg_prio);
```

```
ssize_t mq_receive(mqd_t mqdes, char *msg_ptr,  
                   size_t msg_len, unsigned int *msg_prio)
```


```
int mq_close(mqd_t mqdes)
```

Message Queue – Example

```
send(char* msg){
    mqd_t mq = mq_open("/testqueue", O_RDWR|O_CREAT, 0664, 0);
    if (mq_send(mq, msg, strlen (msg) + 1, 0)<0){
        perror ("MQ Send failure");
        exit (0);
    }
    printf ("MQ Put: %s\n", av [1]);
    return 0;
}
```

```
recieve(){
    mqd_t mq = mq_open("/testqueue",O_RDWR|O_CREAT, 0664,0);
    struct mq_attr attr;
    mq_getattr (mq, &attr); // get attribute
    char *buf = (char*)malloc (attr.mq_msgsize);
    mq_receive(mq, buf, attr.mq_msgsize, NULL);
    printf ("MQ Receive Got: %s\n", buf);
    //clean-up
    mq_close(mq);
    //mq_unlink("/testqueue"); // remove from Kernel
    return 0;
}
```

Shared Memory



- Processes request the segment
- OS maintains the segment – it persists w/o any processes connected
- Processes can map/unmap the segment
- Synchronization is now up to the processes
 - *No send/receive functions, must use “Kernel Semaphores”*

Shared Memory – POSIX functions

- **shm_open**: creates a shared memory segment
- **ftruncate**: sets the size of a shared memory segment
- **mmap**: maps the shared memory object to the process's address space
- **munmap**: unmaps from process's address space
- **shm_unlink**: removes the shared memory segment from the kernel
- **close**: closes the file descriptor associated with the shared memory segment

Shared Memory Example

```
char* my_shm_connect(char* name, int len){
    int fd = shm_open(name, O_RDWR|O_CREAT, 0644 );
    ftruncate(fd, len); //set the length to 1024, the default
    is 0, so this is a necessary step
    char *ptr = (char *) mmap(NULL, len, PROT_READ|PROT_WRITE,
MAP_SHARED, fd, 0); // map
    if (fd < 0){
        perror ("Cannot create shared memory\n");
        exit (0);
    }
    return ptr;
}

void send(char* message){
    char *name = "/testing";
    int len = 1024;
    char* ptr = my_shm_connect (name, len);

    strcpy(ptr, message); // putting data by just copying
    printf ("Put message: %s\n", message);
    close(fd); // close desc, does not remove the segment
    munmap (ptr, len); // this is a bit redundant,
}
```

Shared Memory Example- contd

```
void receive(){
    char    *name    = "/testing";
    int len = 1024;
    char* ptr = my_shm_connect (name, len)

    printf ("Got message: %s\n", ptr);
    shm_unlink (name); //this removes the segment
from Kernel, this is a necessary clean up
    exit(0);
}
```

Kernel Semaphores

- How do we **synchronize processes**?
 - We will again need semaphores, but this time **Kernel Semaphores**
 - They are visible to separate processes who do not share address space
- Operations on Kernel Semaphore:
 - ***sem_open(name, ...)*** to create or connect to a semaphore
 - The name argument must start with a "/"
 - ***sem_close ()*** closes a semaphore
 - It does not destroy it from Kernel
 - ***sem_unlink ()*** removes from Kernel
 - Must be put in the destructor for PA3
 - ***sem_wait ()*** waiting for an event
 - ***sem_post ()*** signaling that an event has occurred
- Find out more from ***sem_overview(7)*** in man pages