



# SIGNALS

Tanzir Ahmed  
CSCE 313 Spring 2021

# Signal: User-Mode Exceptional Flow

- A Signal is a **one-word** message with
  - *Each has a numerical code*
- So far, exceptional control flow features have been usable only by the operating system
  - *Recall handling interrupts and exceptions*
  - *All exception handlers run in **protected (Kernel) mode***
- Would like similar capabilities for **user mode code**
  - *Inter-Process communication to facilitate ‘exceptional control flow’ is subject of today’s discussion*
  - *We will discuss them through “Signals” mechanism*
    - .....but also in a broader context beyond just IPC
- For instance, when we press CTRL-C key we ask the Kernel to send the interrupt signal to the currently running process

# Major Purposes of Signals

1. Obviously, Inter Process Communication
2. A way for humans to interact with programs using the terminal.  
For instance,
  - *Ctrl+C sends a interrupt signal SIGINT*
  - *Ctrl+D sends EOF signal*
  - *Ctrl+Z suspends the process by sending SIGSTP*
3. Most importantly, a mechanism for the Kernel for making processes behave according to specifications.
  - *SIGSEGV is sent on memory exceptions*
  - *SIGILL sent on seeing illegal instructions*
  - *SIGPIPE when on attempts to use a broken pipe*

# List of Signals

Signal	Value	Action	Comment	Terminate and also dump core (for further investigation)
<b>SIGHUP</b>	1	Term	Hangup detected on controlling terminal or death of controlling process	
<b>SIGINT</b>	2	Term	Interrupt from keyboard	
<b>SIGQUIT</b>	3	Core	Quit from keyboard	
<b>SIGILL</b>	4	Core	Illegal Instruction	
<b>SIGABRT</b>	6	Core	Abort signal from <code>abort(3)</code>	
<b>SIGFPE</b>	8	Core	Floating-point exception	
<b>SIGKILL</b>	9	Term	Kill signal	
<b>SIGSEGV</b>	11	Core	Invalid memory reference	
<b>SIGPIPE</b>	13	Term	Broken pipe: write to pipe with no readers; see <code>pipe(7)</code>	
<b>SIGALRM</b>	14	Term	Timer signal from <code>alarm(2)</code>	Only signal ignored by default, unless overridden
<b>SIGTERM</b>	15	Term	Termination signal	
<b>SIGUSR1</b>	30,10,16	Term	User-defined signal 1	
<b>SIGUSR2</b>	31,12,17	Term	User-defined signal 2	
<b>SIGCHLD</b>	20,17,18	Ign	Child stopped or terminated	
<b>SIGCONT</b>	19,18,25	Cont	Continue if stopped	
<b>SIGSTOP</b>	17,19,23	Stop	Stop process	Process scheduler??
<b>SIGTSTP</b>	18,20,24	Stop	Stop typed at terminal	
<b>SIGTTIN</b>	21,21,26	Stop	Terminal input for background process	
<b>SIGTTOU</b>	22,22,27	Stop	Terminal output for background process	

Source of table: <http://man7.org/linux/man-pages/man7/signal.7.html>

Multiple values mean that they are different based on architecture

# What can a Process do about a Signal?

Tell the Kernel what to do with a signal:


1. **Accept Default action.** All signals have a default action  
signal (SIGINT, SIG\_DFL)
2. **Ignore the signal.** Works for most signals  
signal (SIGINT, SIG\_IGN)  
*cannot ignore SIGKILL and SIGSTOP; also unwise to ignore hardware exception signals*
3. **Catch the signal (call a function).** Tell the Kernel to invoke a given function (signal handler) whenever signal occurs.  
signal (SIGINT, foo)

# Signals – Some Points

- The “Action” in the previous page in many cases indicate the “Default Action”
  - You can “**override**” the action for some signals (e.g., **SIGINT**, **SIGUSRx**, **SIGTERM**, **SIGSTP**), while for others (**SIGKILL**, **SIGILL**, **SIGFPE**, **SIGSEGV**, **SIGSTOP**) you **cannot override**
  - “Overriding” is also called “signal handling”, or “catching”
- SIGCHLD is the only signal so far that is **Ignored** by default
  - Most others will kill the process except **SIGCONT** and **SIGSTP**
- Does that mean you can kill any process by sending any signal??
  - NO. You can only signal processes that you created
- **SIGSTP** (suspends a process) and **SIGCONT** (continues/unsuspends a process) are very useful for “**job control**”
  - How to control a process group tied by a shell session?
  - All processes that you run make a group of processes that need special controlling
  - You may suspend a process, run it in the background, move to foreground
  - What happens to standard input/output of a background process?

# Where do Signals come from?

- Today we'll look at Signals in a broader context
  - *IPC is one of the contexts (one process sending to another)*
- Others are facilitated by Users and Kernel
  - *[Users] Signals generated by external Input devices*
  - *[Kernel] Exceptions*



# Where do Signals come from?

(USER) Terminal-generated signals: triggered when user presses certain key on terminal. (e.g. **^C**)

*kill(2) function: Sends any signal to another process.*

*kill(1) command: The command-line interface to kill(2)*

*raise(3) : Sends a signal to itself*

*A user can send signals to only his owned processes*

(Kernel) Exception-generated signals: CPU execution detects condition and notifies Kernel. (e.g. **SIGFPE** divide by 0, **SIGSEGV** invalid memory reference, **SIGILL** when an instruction with illegal opcode is found)

(PROCESSES) Software-condition generated signals: Triggered by software event (e.g. **SIGURG** by out-of-band data on network connection, **SIGPIPE** by broken pipe, **SIGALRM** by timer)





# Generating Signals: `kill(2)` and `raise(3)`

```
#include <signal.h>
```

```
int kill(pid_t pid, int sig);  
/* send signal 'sig' to process 'pid' */
```

```
/* example: send signal SIGUSR1 to process 1234 */  
if (kill(1234, SIGUSR1) == -1)  
    perror("Failed to send SIGUSR1 signal");
```

```
/* example: kill parent process */  
if (kill(getppid(), SIGTERM) == -1)  
    perror("Failed to kill parent");
```

```
#include <signal.h>
```

```
int raise(int sig);  
/* Sends signal 'sig' to itself.  
   Part of ANSI C library! */
```

Raise sends a signal to the executing process  
Kill sends a signal to the specified process

# Signals and the Kernel

- Many of the signals are **generated by** the Kernel in response to events and exceptions received
  - *SIGFPE – FP exception*
  - *SIGILL – Illegal instruction*
  - *SIGSEGV – Segment Violation*
- All others are **routed through** the Kernel, if not originating from the Kernel itself
  - *E.g. SIGHUP (terminal hang), SIGINT (CTRL-C keyboard), SIGTSTP (CTRL-Z), SIGKILL (KILL)*

# Simple Signal Handling: Example



linux2.cse.tamu.edu - PuTTY

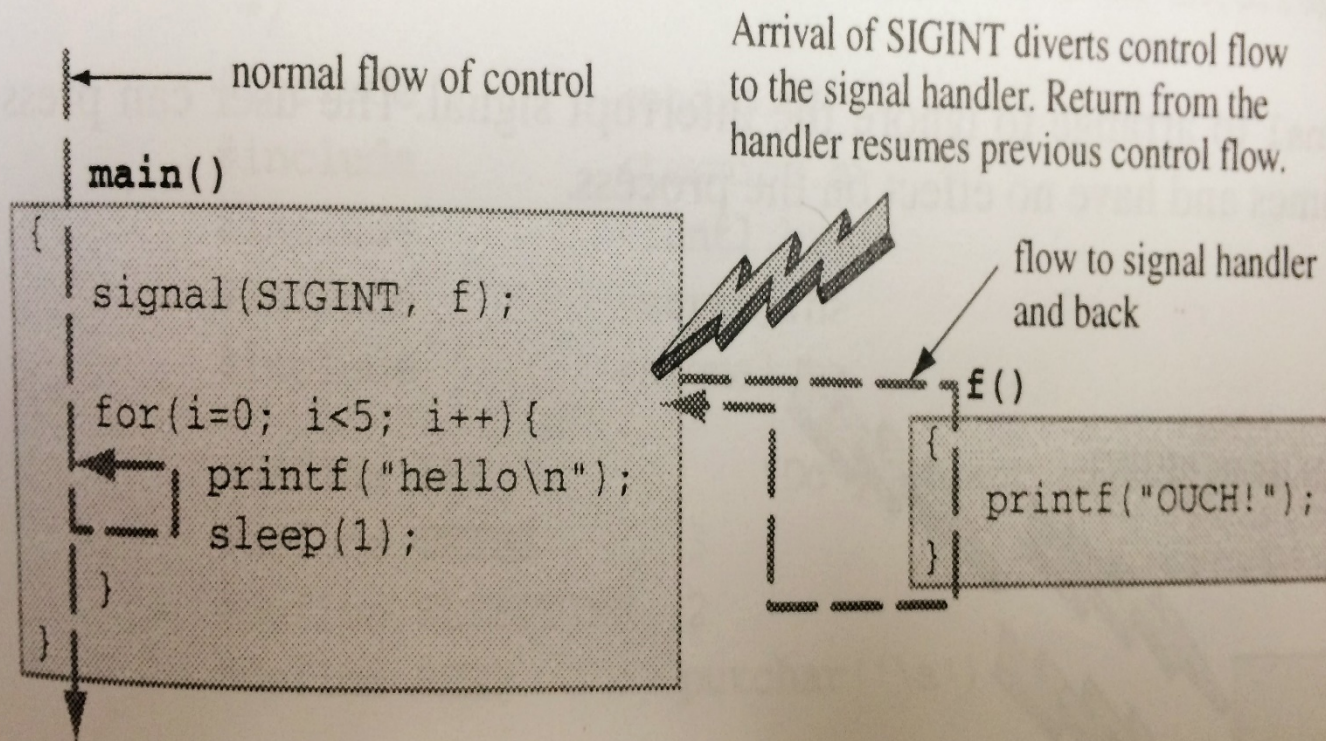
```
:: more sigdemo.c
/* sigdemo.c - shows how a signal handler works
 * credit: Bruce Molay
 */
```

```
#include <stdio.h>
#include <signal.h>
```

```
main()
{
    void f(int); /* declare t
    int i;

    signal(SIGINT, f);
    for (i=0; i<5; i++) {
        printf("hello\n");
        sleep(1);
    }
}
```

```
void f(int signum)
{
    printf("OUCH!\n");
}
```



Taken from: Chapter 6 of "Understanding Unix/Linux Programming" by Bruce Molay

# Example – One Handler for Multiple Signals

```
void sig_usr(int signo) { /*argument is signal number*/
    if      (signo == SIGUSR1) printf("received SIGUSR1\n");
    else if (signo == SIGUSR2) printf("received SIGUSR2\n");
    else    error_dump("received signal %d\n", signo);
    return;
}

int main (void) {
    if (signal(SIGUSR1, sig_usr) == SIG_ERR)
        perror("cannot catch signal SIGUSR1");
    if (signal(SIGUSR2, sig_usr) == SIG_ERR)
        perror("cannot catch signal SIGUSR2");
    for(;;) pause();
}
```

# Reaping Child Processes

- A child process being terminated generates SIGCHLD for the parent process
- The parent can handle this signal by reaping children
- This is asynchronous, i.e., the parent process can do other things instead of continually waiting

```
void my_handler (int sig){
    wait (0);
}

int main (){
    // install handler
    signal (SIGCHLD, my_handler);

    // create 5 child procs,
    // i-th proc sleeps i sec and then dies
    for (int i=1; i<=5; i++){
        int pid = fork ();
        if (pid == 0){
            sleep (i);
            return 0;
        }
    }

    // parent in an infinite loop
    // busy doing something else
    while (true){
        cout << "Relaxing" << endl;
        sleep (1);
    }
}
```

# Signal Blocking

- A process has the option of manually **blocking** the delivery of a signal.
  - *Signal remains blocked until process either (a) unblocks the signal, or (b) changes the action to ignore the signal.*
- Blocking is automatic for a signal type when the same signal is being handled
  - *First, SIGINT received causes the SIGINT handler (if any) to be called*
  - *Also causes SIGINT to be blocked for the process - If another SIGINT occurs during Handler execution, it is recorded in the pending bit vector, but not delivered*
  - *When handler returns, signals of that type can start to be delivered again*
- The following function is used to block a signal:
  - ***sigprocmask*** (*manipulate the set of blocked signals for the process*)

# Properly Reaping Child Processes Under Simultaneous Termination

- A child process being terminated generates SIGCHLD for the parent process
- The parent can handle this signal by reaping children
- This is asynchronous, i.e., the parent process can do other things instead of continually waiting

```
void my_handler (int sig){
    while (waitpid (-1, 0, 0) != -1)
        ;
}

int main (){
    // install handler
    signal (SIGCHLD, my_handler);

    // create 5 child procs,
    // i-th proc sleeps i sec and then dies
    for (int i=1; i<=5; i++){
        int pid = fork ();
        if (pid == 0){
            sleep (5); //all die
            return 0;
        }
    }

    // parent in an infinite loop
    // busy doing something else
    while (true){
        cout << "Relaxing" << endl;
        sleep (1);
    }
}
```