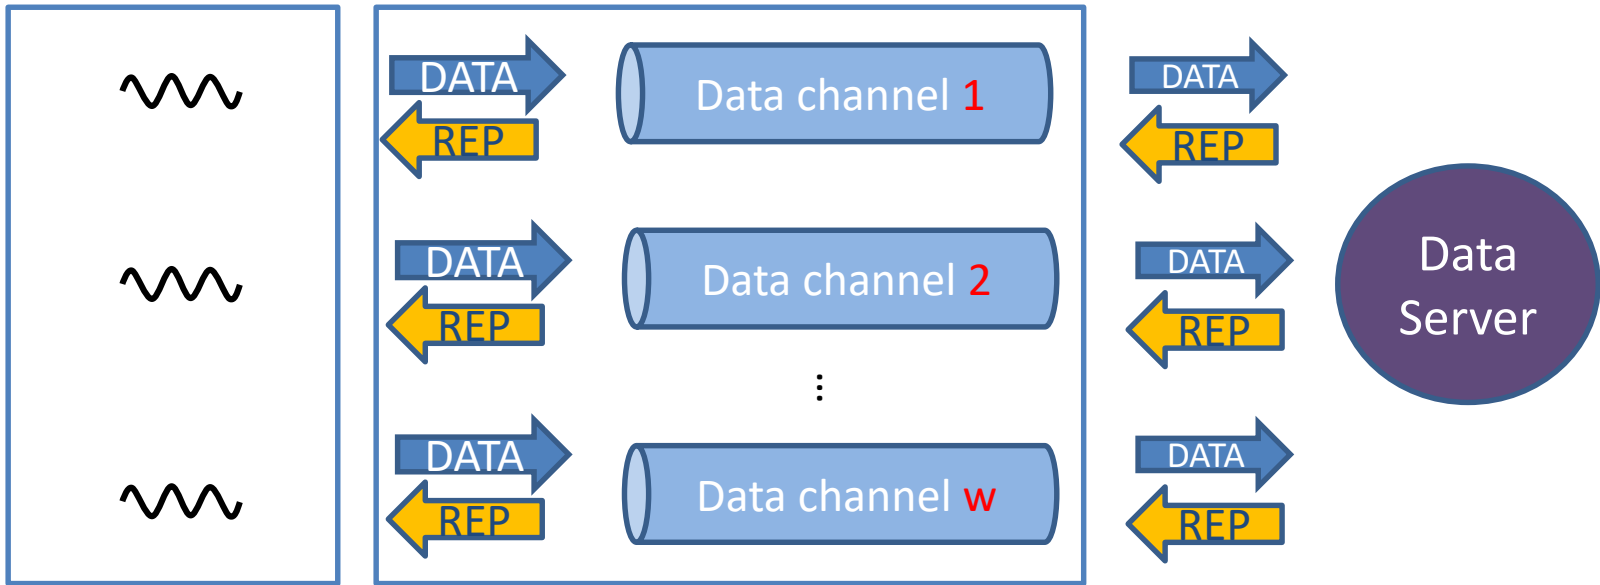


CSCE 313

PA5: I/O Multiplexing

PA4 Review

w worker threads



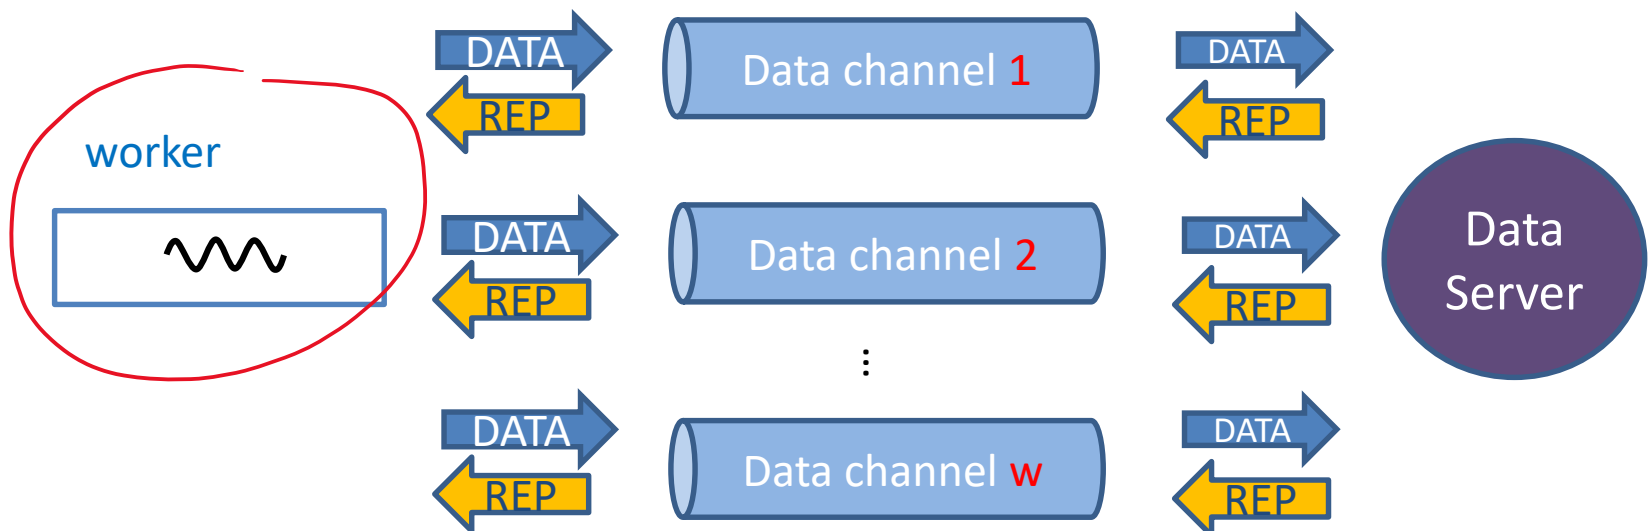
- Each worker spends most of the time waiting for server responses (inter-process delay)
 - Assume 10 sec communication delay
 - 1 worker can do 0.1 req/sec, 10 workers can do 1 req/sec

PA4 Review

- Each worker watches one data channel for response
 - Blocking I/O: Wait till data is ready
 - Low CPU utilization as most workers are waiting
- Workers are constantly switched on and off CPU
 - High context switching cost
 - Given x CPU cores, is there any reason to have more than x workers?


PA5

- PA5 Method: Use one single worker thread to deal with the I/O for all data channels
 - Avoid context switching
 - Non-blocking I/O: After issuing a request, switches to other tasks without waiting for response
 - When data from data channel i is ready, the client detects that by waiting on `epoll()`

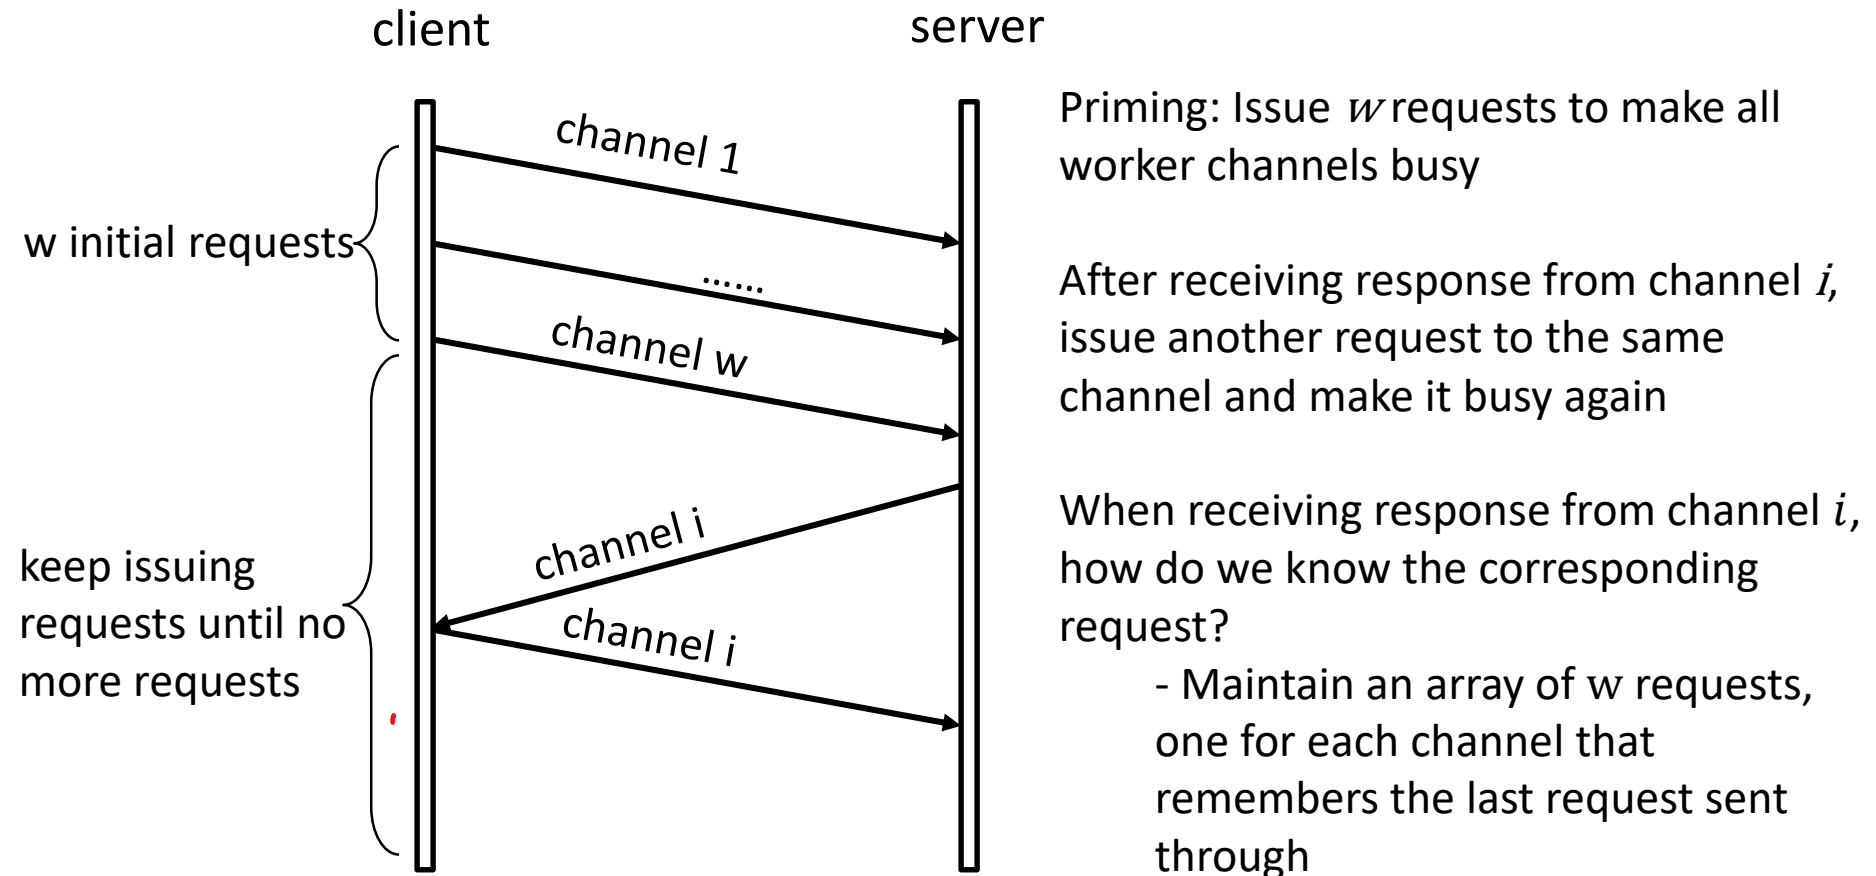


What We Did in PA4

- In PA4, we called `cwrite()` immediately followed by a `cread()` – this made a thread always until data comes back from server
- In PA5, we will change that by issuing a bunch of `cwrite()`'s and then a bunch of `cread()`'s

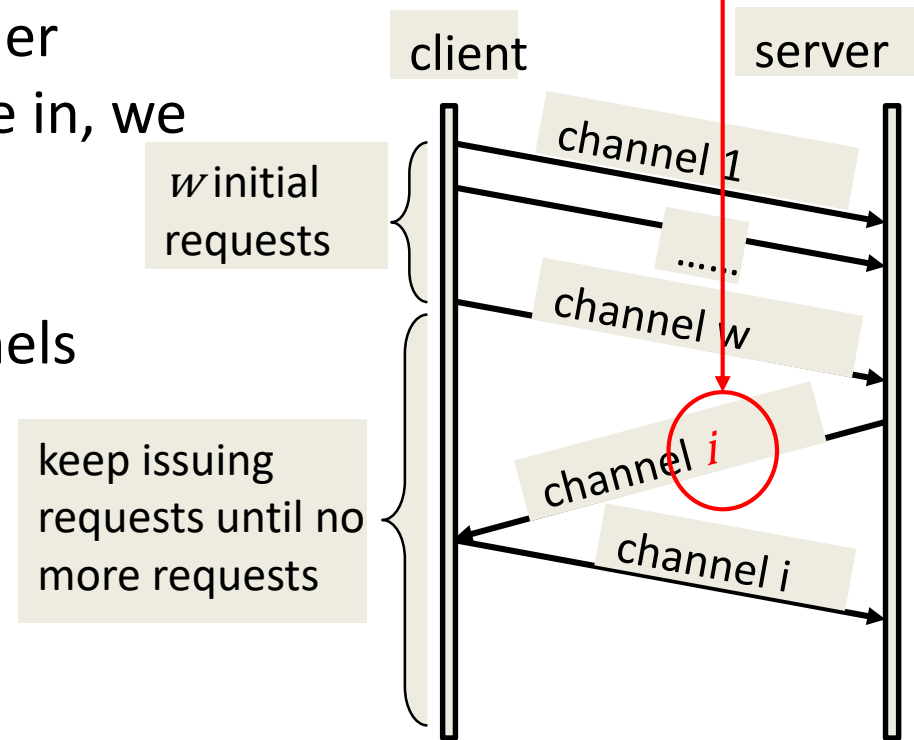


PA5



The Big Question

- How do we know the perfect order in which to receive responses?
 - You cannot expect the responses will arrive in the same order as the corresponding requests were sent
 - What is the value of i in the picture?
- If we wait for responses in an order different from the one they arrive in, we would be waiting unnecessarily
 - Leads to inefficiency
- We will still “block” for the channels as a group, not individually
- The function to use is `epoll()`



Epoll - Usage

- Use `epoll()` to listen to `w` worker channels
 - It is a collection of file descriptors from our request channels
 - `epoll_create1()` creates the empty list
 - `epoll_ctl()` lets us add file descriptors (i.e., `FIFORequestChannel::rfd`'s, no need for the `wfd`'s) from each channel
 - `epoll_wait()` waits until one or more in the list has some activity, i.e., data available to read
 - The best place to look is `epoll(7)` linux man pages

```
int epoll_create1(int flags);
```

```
int epoll_ctl(int epfd, int op, int fd,  
              struct epoll_event *event);
```

```
int epoll_wait(int epfd, struct epoll_event *events,  
               int maxevents, int timeout);
```