

Name:.....Chaoyang Zhu..... UIN.....930001840...
 Total Points: -----/100

1. [5 pts] While implementing the process state diagram, what is the problem of having only 1 queue for all blocked processes waiting for all events? What is the solution to this problem? Describe with an example.

The efficiency of dealing with different types of events is low. It does not make sense to make a file request from disk and URL request wait in the same queue. We can solve this problem by separating different kinds of events into different queues.

2. [25 pts] Assume the following processes A, B, C are loaded in memory of a system that uses both multiprogramming(MP) and timesharing (TS) techniques. These processes have 15, 7, and 13 instructions respectively. Also assume that the dispatcher lives at address 100 in memory and spans 4 instructions (i.e., 100-103). The time quantum is long enough for exactly **5 instructions**.

The following table shows only instruction addresses in the memory with I/O requests labelled, along with the duration of these I/O operations in terms of CPU instructions. Although I/O operations do not take CPU instructions, the duration means that the I/O operations will finish by the time the corresponding number of CPU instructions execute.

Please draw one possible trace of these 3 processes running together in the CPU. Use page 14-15 of Lecture 3 for reference. You may skip the first invocation of the dispatcher to decide the first process to run in the CPU.

Process A	Process B	Process C
5000	8000	12000
5001	8001	12001(I/O, takes 3 ins.)
5002	8002	12002
5003	8003 (I/O, takes 7 ins.)	12003
5004 (I/O, takes 6 ins.)	8004	12004
5005 (I/O, takes 4 ins.)	8005	12005
5006	8006	12006(I/O, takes 1 ins.)
5007		12007
5008(I/O, takes 5 ins.)		12008(I/O, takes 2 ins.)
5009		12009
5010		12010
5011		12011
5012		12012

5013 5014		
--------------	--	--

My answer:

1 5000

2 5001

3 5002

4 5003

5 5004

----- Timeout and I/O Request (6 ins)

6 100

7 101

8 102

9 103

10 8000

11 8001

12 8002

13 8003

----- I/O Request (7 ins)

14 100

15 101

16 102

17 103

18 12000

19 12001

----- I/O Request (3 ins)

20 100

21 101

22 102

23 103

24 5005

----- I/O Request (4 ins)

25 100

26 101

27 102

28 103

29 8004

30 8005

31 8006

32 100

33 101

34 102

35 103

36 12002

37 12003

38 12004

39 12005

40 12006

----- I/O Request (1 ins)

41 100

42 101

43 102

44 103

45 5006

46 5007

47 5008

----- I/O Request (5 ins)

48 100

49 101

50 102

51 103

52 12007

53 12008

----- I/O Request (2 ins)

54 100

55 101

56 102

57 103

58 5009

59 5010

60 5011

61 5012

62 5013

----- Timeout

63 100

64 101

65 102

66 103

67 12009

68 12010

69 12011

70 12012

----- I/O Request

71 100

72 101

73 102

74 103

75 5014

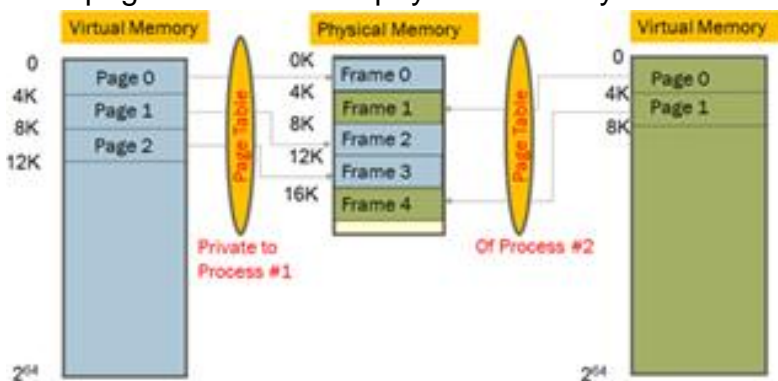
3. [5 pts] What is the difference between the "New" state and the "Ready to Run" state in the process state diagram?

The "New" means the PCB for the process was created, but the executable image has not been loaded into the memory. While "Ready to Run" means that the executable image has been loaded and the program can be executed.

4. [5 pts] Is a transition from the "Blocked" state to directly to the "Exit" state possible in the process state diagram? How?

Yes, it is possible. For example, parent process killed the child while it was waiting.


5. [10 pts] Assume that the following physical memory is full with already allocated 5 pages as shown below (i.e., it is 20KB in capacity). Describe what happens if process 2 wants to allocate and use another page. What changes in the page tables and the physical memory?



When process 2 wants to allocate and use another page, the Page Fault happens. The fault-handler allocates the required page into the physical memory. The Least Recently Used frame will be evicted to the disk. Meanwhile, the corresponding page table will be updated.

6. [25 pts] Draw the Descriptor Table (DT) and File Table (FT) for each possible process just after executing the lines 2, 5 and 10 and explain what you see in the standard output and **file.txt** along the way. FT entries should contain the cursor and the reference count


```
1.int fd=open("file.txt", O_CREAT|O_RDWR);
2.write (fd, "Hello world", 6);
3.if (!fork()){
4.  dup2 (fd, 1); //redirect stdout
5.  cout << "Mars, I am here!!" << endl;
6.}else{
7.  wait(0);
8.  lseek (fd, 0, SEEK_SET);
9.  cout << "Mars is great" << endl;
10. write (fd, "Hola  ", 5); //draw tables
11. close (fd);
12.}
```



file.txt : "Hello "

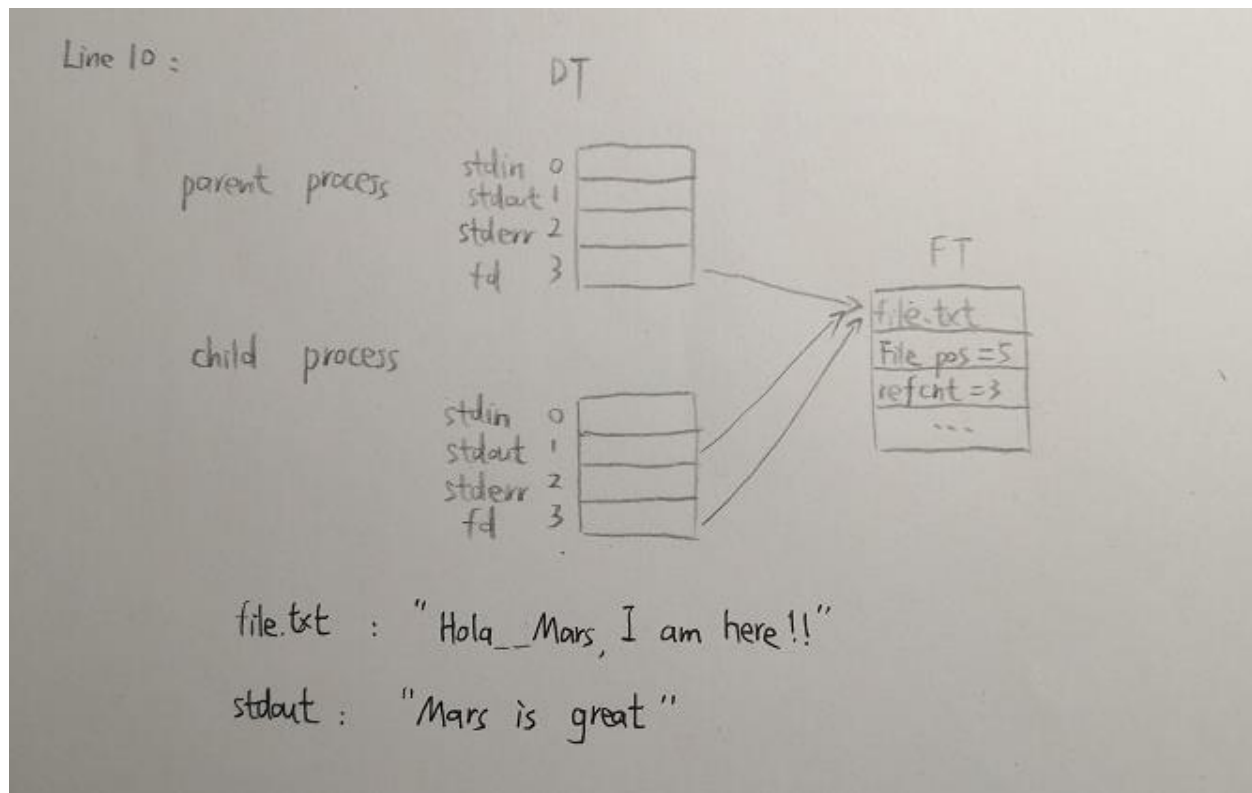
stdout : nothing

Line 5:



file.txt : "Hello Mars, I am here !!"

in child process, stdout points to file.txt



7. [25 pts] Please explain the output of the following program using a process tree diagram as shown in Lecture 4 page 23. Assuming the main process's ID is 1000, explain the order of process creation (in fact, the number itself should show what order the processes are created). Also explain what you see in the standard output. Are different outputs possible for this program? Why or why not?

```
for (int i=0; i<4; i++){
    int cid = fork ();
    if (i < 2)
        wait (0);
    cout << "ID=" << getpid () << endl;
}
```

If there is a wait between two forks, all the child processes of the former fork will print out their pid numbers before the latter fork.

If there is no wait between two forks, the order of pid numbers being printed out in different layers of fork is undetermined. Since the system scheduler algorithm is different.

Different outputs are possible for this program in different systems because of the different scheduler algorithm.

